

3.1_stochastic_gradient_descent_v3

March 25, 2022

Linear regression 1D: Training Two Parameter Stochastic Gradient Descent (SGD)

Objective

How to use SGD(Stochastic Gradient Descent) to train the model.

Table of Contents

In this Lab, you will practice training a model by using Stochastic Gradient descent.

Make Some Data

Create the Model and Cost Function (Total Loss)

Train the Model:Batch Gradient Descent

Train the Model:Stochastic gradient descent

Train the Model:Stochastic gradient descent with Data Loader

Estimated Time Needed: 30 min

Preparation

We'll need the following libraries:

```
[1]: # These are the libraries we are going to use in the lab.

import torch
import matplotlib.pyplot as plt
import numpy as np

from mpl_toolkits import mplot3d
```

The class `plot_error_surfaces` is just to help you visualize the data space and the parameter space during training and has nothing to do with PyTorch.

```
[2]: # The class for plot the diagram

class plot_error_surfaces(object):

    # Constructor
    def __init__(self, w_range, b_range, X, Y, n_samples = 30, go = True):
        W = np.linspace(-w_range, w_range, n_samples)
```

```

B = np.linspace(-b_range, b_range, n_samples)
w, b = np.meshgrid(W, B)
Z = np.zeros((30, 30))
count1 = 0
self.y = Y.numpy()
self.x = X.numpy()
for w1, b1 in zip(w, b):
    count2 = 0
    for w2, b2 in zip(w1, b1):
        Z[count1, count2] = np.mean((self.y - w2 * self.x + b2) ** 2)
        count2 += 1
    count1 += 1
self.Z = Z
self.w = w
self.b = b
self.W = []
self.B = []
self.LOSS = []
self.n = 0
if go == True:
    plt.figure()
    plt.figure(figsize = (7.5, 5))
    plt.axes(projection = '3d').plot_surface(self.w, self.b, self.Z,
↪rstride = 1, cstride = 1, cmap = 'viridis', edgecolor = 'none')
    plt.title('Loss Surface')
    plt.xlabel('w')
    plt.ylabel('b')
    plt.show()
    plt.figure()
    plt.title('Loss Surface Contour')
    plt.xlabel('w')
    plt.ylabel('b')
    plt.contour(self.w, self.b, self.Z)
    plt.show()

# Setter
def set_para_loss(self, W, B, loss):
    self.n = self.n + 1
    self.W.append(W)
    self.B.append(B)
    self.LOSS.append(loss)

# Plot diagram
def final_plot(self):
    ax = plt.axes(projection = '3d')
    ax.plot_wireframe(self.w, self.b, self.Z)

```

```

        ax.scatter(self.W, self.B, self.LOSS, c = 'r', marker = 'x', s = 200,
↪alpha = 1)
        plt.figure()
        plt.contour(self.w, self.b, self.Z)
        plt.scatter(self.W, self.B, c = 'r', marker = 'x')
        plt.xlabel('w')
        plt.ylabel('b')
        plt.show()

        # Plot diagram
        def plot_ps(self):
            plt.subplot(121)
            plt.ylim
            plt.plot(self.x, self.y, 'ro', label = "training points")
            plt.plot(self.x, self.W[-1] * self.x + self.B[-1], label = "estimated_
↪line")
            plt.xlabel('x')
            plt.ylabel('y')
            plt.ylim((-10, 15))
            plt.title('Data Space Iteration: ' + str(self.n))
            plt.subplot(122)
            plt.contour(self.w, self.b, self.Z)
            plt.scatter(self.W, self.B, c = 'r', marker = 'x')
            plt.title('Loss Surface Contour Iteration' + str(self.n))
            plt.xlabel('w')
            plt.ylabel('b')
            plt.show()

```

Make Some Data

Set random seed:

```

[3]: # Set random seed

torch.manual_seed(1)

```

```

[3]: <torch._C.Generator at 0x7feb7c025370>

```

Generate values from -3 to 3 that create a line with a slope of 1 and a bias of -1. This is the line that you need to estimate. Add some noise to the data:

```

[4]: # Setup the actual data and simulated data


X = torch.arange(-3, 3, 0.1).view(-1, 1)
f = 1 * X - 1
Y = f + 0.1 * torch.randn(X.size())

```

Plot the results:

```
[5]: # Plot out the data dots and line

plt.plot(X.numpy(), Y.numpy(), 'rx', label = 'y')
plt.plot(X.numpy(), f.numpy(), label = 'f')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
plt.show()
```



Create the Model and Cost Function (Total Loss)

Define the forward function:

```
[6]: # Define the forward function
```

```
def forward(x):
    return w * x + b
```

Define the cost or criterion function (MSE):

```
[7]: # Define the MSE Loss function
```


```
def criterion(yhat, y):
    return torch.mean((yhat - y) ** 2)
```


Create a `plot_error_surfaces` object to visualize the data space and the parameter space during

training:

```
[8]: # Create plot_error_surfaces for viewing the data  
get_surface = plot_error_surfaces(15, 13, X, Y, 30)
```

<Figure size 432x288 with 0 Axes>





Train the Model: Batch Gradient Descent

Create model parameters w , b by setting the argument `requires_grad` to `True` because the system must learn it.

```
[9]: # Define the parameters  $w$ ,  $b$  for  $y = wx + b$ 

w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
```

Set the learning rate to 0.1 and create an empty list `LOSS` for storing the loss for each iteration.

```
[10]: # Define learning rate and create an empty list for containing the loss for
      ↪ each iteration.

lr = 0.1
LOSS_BGD = []
```

Define `train_model` function for train the model.

```
[11]: # The function for training the model

def train_model(iter):

    # Loop
```

```

for epoch in range(iter):

    # make a prediction
    Yhat = forward(X)

    # calculate the loss
    loss = criterion(Yhat, Y)

    # Section for plotting
    get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.
↪tolist())
    get_surface.plot_ps()

    # store the loss in the list LOSS_BGD
    LOSS_BGD.append(loss)

    # backward pass: compute gradient of the loss with respect to all the ↵
↪learnable parameters
    loss.backward()

    # update parameters slope and bias
    w.data = w.data - lr * w.grad.data
    b.data = b.data - lr * b.grad.data

    # zero the gradients before running the backward pass
    w.grad.data.zero_()
    b.grad.data.zero_()

```



Run 10 epochs of batch gradient descent: bug data space is 1 iteration ahead of parameter space.



```



[12]: # Train the model with 10 iterations



train_model(10)



```











Train the Model: Stochastic Gradient Descent

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

```
[13]: # Create plot_error_surfaces for viewing the data

get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Define `train_model_SGD` function for training the model.

```
[14]: # The function for training the model

LOSS_SGD = []
w = torch.tensor(-15.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)

def train_model_SGD(iter):

    # Loop
    for epoch in range(iter):

        # SGD is an approximation of our true total loss/cost, in this line of
        ↪code we calculate our true loss/cost and store it
        Yhat = forward(X)

        # store the loss
        LOSS_SGD.append(criterion(Yhat, Y).tolist())

        for x, y in zip(X, Y):

            # make a prediction
            yhat = forward(x)

            # calculate the loss
            loss = criterion(yhat, y)

            # Section for plotting
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.
            ↪tolist())

            # backward pass: compute gradient of the loss with respect to all
            ↪the learnable parameters
            loss.backward()

            # update parameters slope and bias
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
```

```

# zero the gradients before running the backward pass
w.grad.data.zero_()
b.grad.data.zero_()


#plot surface and data space after each epoch
get_surface.plot_ps()



```



Run 10 epochs of stochastic gradient descent: bug data space is 1 iteration ahead of parameter space.

```
[15]: # Train the model with 10 iterations
```


```
train_model_SGD(10)
```









Data Space Iteration: 360 Loss Surface Contour Iteration360




Data Space Iteration: 420 Loss Surface Contour Iteration420




Data Space Iteration: 480 Loss Surface Contour Iteration 480



Data Space Iteration: 540 Loss Surface Contour Iteration 540






Compare the loss of both batch gradient descent as SGD.

```
[20]: # Plot out the LOSS_BGD and LOSS_SGD

plt.plot(LOSS_BGD,label = "Batch Gradient Descent")
plt.plot(LOSS_SGD,label = "Stochastic Gradient Descent")
plt.xlabel('epoch')
plt.ylabel('Cost/ total loss')
plt.legend()
plt.show()
```



SGD with Dataset DataLoader

Import the module for building a dataset class:

```
[21]: # Import the library for DataLoader

from torch.utils.data import Dataset, DataLoader
```

Create a dataset class:

```
[22]: # Dataset Class

class Data(Dataset):

    # Constructor
    def __init__(self):
        self.x = torch.arange(-3, 3, 0.1).view(-1, 1)
        self.y = 1 * self.x - 1
        self.len = self.x.shape[0]

    # Getter
    def __getitem__(self, index):
        return self.x[index], self.y[index]

    # Return the length
    def __len__(self):
```

```
return self.len
```

Create a dataset object and check the length of the dataset.

```
[23]: # Create the dataset and check the length

dataset = Data()
print("The length of dataset: ", len(dataset))
```

The length of dataset: 60

Obtain the first training point:

```
[24]: # Print the first point

x, y = dataset[0]
print("(", x, ", ", y, ")")
```

(tensor([-3.]) , tensor([-4.]))

Similarly, obtain the first three training points:

```
[25]: # Print the first 3 point

x, y = dataset[0:3]
print("The first 3 x: ", x)
print("The first 3 y: ", y)
```

The first 3 x: tensor([[-3.0000],
[-2.9000],
[-2.8000]])

The first 3 y: tensor([[-4.0000],
[-3.9000],
[-3.8000]])

Create a `plot_error_surfaces` object to visualize the data space and the parameter space during training:

```
[26]: # Create plot_error_surfaces for viewing the data

get_surface = plot_error_surfaces(15, 13, X, Y, 30, go = False)
```

Create a `DataLoader` object by using the constructor:

```
[27]: # Create DataLoader

trainloader = DataLoader(dataset = dataset, batch_size = 1)
```

Define `train_model_DataLoader` function for training the model.

```

[28]: # The function for training the model

w = torch.tensor(-15.0,requires_grad=True)
b = torch.tensor(-10.0,requires_grad=True)
LOSS_Loader = []

def train_model_DataLoader(epochs):

    # Loop
    for epoch in range(epochs):

        # SGD is an approximation of our true total loss/cost, in this line of
        ↪code we calculate our true loss/cost and store it
        Yhat = forward(X)

        # store the loss
        LOSS_Loader.append(criterion(Yhat, Y).tolist())

        for x, y in trainloader:

            # make a prediction
            yhat = forward(x)

            # calculate the loss
            loss = criterion(yhat, y)

            # Section for plotting
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.
            ↪tolist())

            # Backward pass: compute gradient of the loss with respect to all
            ↪the learnable parameters
            loss.backward()

            # Update parameters slope
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data

            # Clear gradients
            w.grad.data.zero_()
            b.grad.data.zero_()


            #plot surface and data space after each epoch
            get_surface.plot_ps()

```


Run 10 epochs of stochastic gradient descent: bug data space is 1 iteration ahead of parameter space.

```
[29]: # Run 10 iterations
```


```
train_model_DataLoader(10)
```




Data Space Iteration: 120 Loss Surface Contour Iteration 120




Data Space Iteration: 180 Loss Surface Contour Iteration 180





Data Space Iteration: 240 Loss Surface Contour Iteration 240




Data Space Iteration: 300 Loss Surface Contour Iteration 300







Data Space Iteration: 480 Loss Surface Contour Iteration 480



Data Space Iteration: 540 Loss Surface Contour Iteration 540






Compare the loss of both batch gradient decent as SGD. Note that SGD converges to a minimum faster, that is, it decreases faster.

```
[30]: # Plot the LOSS_BGD and LOSS_Loader

plt.plot(LOSS_BGD,label="Batch Gradient Descent")
plt.plot(LOSS_Loader,label="Stochastic Gradient Descent with DataLoader")
plt.xlabel('epoch')
plt.ylabel('Cost/ total loss')
plt.legend()
plt.show()
```





Practice



For practice, try to use SGD with DataLoader to train model with 10 iterations. Store the total loss in LOSS. We are going to use it in the next question.



```
[31]: # Practice: Use SGD with trainloader to train model and store the total loss in LOSS
      ↪ LOSS



LOSS = []
w = torch.tensor(-12.0, requires_grad = True)
b = torch.tensor(-10.0, requires_grad = True)
def my_train_model(epochs):
    for epoch in range(epochs):
        Yhat = forward(X)
        LOSS.append(criterion(Yhat, X))
        for x, y in trainloader:
            yhat = forward(x)
            loss = criterion(yhat, y)
            get_surface.set_para_loss(w.data.tolist(), b.data.tolist(), loss.
            ↪ tolist())
            loss.backward()
            w.data = w.data - lr * w.grad.data
            b.data = b.data - lr * b.grad.data
            w.grad.data.zero_()
            b.grad.data.zero_()
```



```
get_surface.plot_ps()  
my_train_model(10)
```












Double-click [here](#) for the solution.

Plot the total loss

```
[34]: # Practice: Plot the total loss using LOSS

# Type your code here
for i in range(len(LOSS)):
    LOSS[i] = LOSS[i].item()
plt.plot(LOSS, label = "Stochastic Gradient Descent")
plt.xlabel('iteration')
plt.ylabel('Cost/ total loss')
plt.legend()
plt.show()
```



Double-click [here](#) for the solution.

About the Authors:

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: Michelle Carey, Mavis Zhou

Thanks to: Andrew Kin ,Alessandro Barboza

0.1 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-09-23	2.0	Shubham	Migrated Lab to Markdown and added to course repo in GitLab

##

© IBM Corporation 2020. All rights reserved.